



Fireflower detection

Javae Elliott and Joseph Strizhak

Problem Motivation



Air Patrols

Over vast amounts of suspected areas



Thermal Imaging

Detect hot spots & residual fires



Lookout Towers

Manned towers with extensive view



Local Reports

By general public, commercial & recreational pilots



Weather Monitoring

Lightning strikes, Moisture content, probability of wildfire

Executive Summary



Goal

Vision-based
model to provide
wildfire detection



Solution

Provide a CNN for
object detection &
localization



Benefit

Cost, Assistance to
already-existing
methods of
detection



Technical Problems

Data Relevance

Most accessible
were not specific to
wildfires



Tensorflow Datasets

TFDS, AutoKeras, HParams,
Tensorboard



Data Quality

Dataset size, image
quality, curation

Performance

Model overfitting, metrics
like f1 score 'undoable'

Describe the feature and the current behavior/state.

Currently, F1-score cannot be meaningfully used as a metric in keras neural network models, because keras will call F1-score at each batch step at validation, which results in too small values. Therefore, F1-score was removed from keras, see [keras-team/keras#5794](https://keras-team.github.io/keras/#5794), where also some quick solution is proposed. Tfa's F1-score exhibits exactly the same problem when used with keras.

Relevant information



BIOLOGICAL STRATEGY



The Beak That Inspired a Bullet Train Kingfishers



Our Concept

In the same line of thinking, we thought of ways to get around the problem of fire image data quality.



Our Concept

Transfer learn a model trained on the  tf_flowers dataset onto our best curated* fire dataset

*removed non-natural objects, people. Kept only fire region of searched images. Balanced

Challenges & Background



HEAVILY
undocumented



TFDS

Use TFDS so as to
avoid RAM issues



AutoKeras
Tensorboard
HParams



Fire Dataset

1900 RGB images
950 per class



tf_flowers

3670 RGB images of
6 classes


```
clf = ak.ImageClassifier(
    overwrite=True,
    max_trials=10,
    num_classes=5,
    #tuner=>AutoKeras' Default tuner
    loss='sparse_categorical_crossentropy',
    directory='/content/',
    seed=seed,
    objective='val_accuracy'
    #metric<-default accuracy
)

clf.fit(
    fl_train_data,
    batch_size=BATCH_SIZE,#put in to have autokeras change it if too much memory
    validation_split=0.2,
    epochs=5
)

print(clf.evaluate(fl_test_data))

Trial 10 Complete [00h 09m 28s]
val_accuracy: 0.9137324094772339

Best val_accuracy So Far: 0.9260563254356384
Total elapsed time: 01h 08m 22s
```


```
def train_test_model(hparams):
    model = tf.keras.Sequential()
    for layer in autokeras_model.layers[:-2]: # Skip first and last layer
        model.add(layer)
    model.add(tf.keras.layers.Dropout(hparams[HP_DROPOUT]))
    model.add(tf.keras.layers.Dense(fire_num_classes, activation='softmax'))

    model.compile(
        optimizer = tf.keras.optimizers.Adam(learning_rate=hparams[HP_LR]),
        loss = tf.keras.losses.SparseCategoricalCrossentropy(),
        metrics=['accuracy'])

    history = model.fit(train_fire,
                        epochs=init_epochs)
    _,accuracy = model.evaluate(valid_fire)
    return accuracy

def run(run_dir, hparams):
    with tf.summary.create_file_writer(run_dir).as_default():
        hp.hparams(hparams) # record the values used in this trial
        accuracy = train_test_model(hparams)
        tf.summary.scalar(METRIC_ACCURACY, accuracy, step=1)
```

Approach

We attempt a neural architecture search using AutoKeras on the  tf_flowers dataset.

Then, using feature extraction, transfer learn that model onto the fire dataset, using Tensorboard & HParams to optimize performance.



Hyperparameters

learning_rate

Min

-infinity

Max

+infinity

dropout

Min

-infinity

Max

+infinity

optimizer

adam

Metrics

accuracy

Min


-infinity

Max

+infinity

accuracy

tag: accuracy



| Name | Smoothed | Value | Step | Time |
|-------|----------|--------|------|----------------------|
| run-0 | 0.9757 | 0.9757 | 1 | Sun Dec 18, 14:17:54 |
| run-1 | 0.9826 | 0.9826 | 1 | Sun Dec 18, 14:18:56 |
| run-2 | 0.9826 | 0.9826 | 1 | Sun Dec 18, 14:20:06 |
| run-3 | 0.9861 | 0.9861 | 1 | Sun Dec 18, 14:21:08 |
| run-4 | 0.9861 | 0.9861 | 1 | Sun Dec 18, 14:22:12 |
| run-6 | 0.9757 | 0.9757 | 1 | Sun Dec 18, 14:24:16 |
| run-7 | 0.9792 | 0.9792 | 1 | Sun Dec 18, 14:25:18 |
| run-8 | 0.9792 | 0.9792 | 1 | Sun Dec 18, 14:26:19 |

TABLE VIEW

PARALLEL COORDINATES VIEW

SCATTER PLOT MATRIX VIEW

| Trial ID | Show Metrics | learning_rate | dropout | optimizer | accuracy |
|------------------|--------------------------|---------------|---------|-----------|----------|
| 018c7e51f03f2... | <input type="checkbox"/> | 0.0064672 | 0.10000 | | 1.0000 |
| 3338a144cd8b4... | <input type="checkbox"/> | 0.010000 | 0.20000 | | 0.98611 |
| 39df7384dab67... | <input type="checkbox"/> | 0.010000 | 0.10000 | | 0.98611 |
| 75ed77221f766... | <input type="checkbox"/> | 0.0064672 | 0.18260 | | 0.97917 |
| 87e2e12931d42... | <input type="checkbox"/> | 0.0010000 | 0.19583 | | 0.98611 |
| 99b73fa562c44... | <input type="checkbox"/> | 0.010000 | 0.10379 | | 0.98958 |
| 9ca752c345ae4... | <input type="checkbox"/> | 0.0010000 | 0.10000 | adam | 0.97569 |
| a17636c40bcf9... | <input type="checkbox"/> | 0.0010000 | 0.20000 | | 0.97917 |
| b138e04f2ec2c... | <input type="checkbox"/> | 0.0064672 | 0.20000 | | 0.98611 |

| Layer (type) | Output Shape | Param # |
|---|--------------------------|----------|
| cast_to_float32 (CastToFloat32) | (None, 250, 250, 3) | 0 |
| normalization (Normalization) | (None, 250, 250, 3) | 7 |
| random_flip (RandomFlip) | (None, 250, 250, 3) | 0 |
| random_rotation (RandomRotation) | (None, 250, 250, 3) | 0 |
| efficientnetb7 (Functional) | (None, None, None, 2560) | 64097687 |
| global_average_pooling2d (GlobalAveragePooling2D) | (None, None, None, 2560) | 0 |
| dropout_2 (Dropout) | (None, 2560) | 0 |
| dense_1 (Dense) | (None, 2) | 5122 |
| Total params: 64,102,816 | | |
| Trainable params: 5,122 | | |
| Non-trainable params: 64,097,694 | | |

Solution & Implementation

AutoKeras found a model (EfficientNet pre-trained on ImageNet) that yielded top validation accuracy.

Tensorboard found minimal gain in feature extraction between dropout & learning rate.

We end with (on test set)

Accuracy: 0.955263 Loss: 0.128732

Localization

Transfer learning

- Pre-trained model from Tensorflow Object Detection API

Bounding boxes for fires

- Used label studio for picture labeling

Re-training/fine-tuning

- Retrained the Tensorflow model for fire detection



Possibilities & Improvements


Video, not Image

Object detection models trained on videos are more robust



BoundingBox Space

Some images are engulfed in flame, its effect on localization would be harmful



Tensorboard

Tensorboard has a lot more visualizational range that would've been ideal to explore

Precision

Due to API issues, we can't exactly measure some relevant metrics for our classifier.



Thanks!

Javae Elliott and Joseph Strizhak

[GitHub Link](#)

Fireflower detection

CREDITS: This presentation template was
created by **Slidesgo**, including icons by **Flaticon**,
and infographics & images by **Freepik**